

# TCP Client-Server Stack Implementation



COSC 414 - Professor Trivedi

Department of Computer Science: Operating Systems

Ihab Ashkar

Chibueze Eburuoh

Xannia Simpson

# Content

1 Project Overview



2 System Architecture



3 Key Implementation Details



4 Server Implementation



5 Client Implementation



6 Build Process



7 OS IPC Breakdown



8 Future Enhancements



# Project Overview

This project implements a simplified client–server communication model to demonstrate how data moves through a basic network stack in a UNIX-based operating system. Using TCP sockets, the system shows how processes exchange messages through network-based inter-process communication (IPC). Although minimal by design, the implementation highlights key networking concepts such as socket creation, binding, listening, connecting, and data transmission—mirroring the essential behavior of real OS-level network layers.



# System Architecture

## Client to Server Communication

1. Client creates socket and requests connection.
2. Server listens on port 8080 using socket(), bind(), and listen().
3. Server accepts client and creates dedicated connection socket.
4. Client sends message, server processes, and server responds.
5. Both sides close connection.

## Files

- server.cpp handles socket setup, client handling, and the message response
- client.cpp creates the connection, sends the connection request, and prints the server reply.
- Our Makefile simply automates compilation for both files, rather than manually compiling each.

# Key Implementation Details

- **Server Setup:**
  - Create TCP socket
  - Bind to a port
  - Listen for a connection request
- **Client Handling:**
  - Accept a client socket
  - Use **recv()** (receive) and **send()** for communication
  - Loop until client disconnects
- **Client Operations:**
  - Create socket
  - Connect to server using **connect()**
  - Send message and wait for server reply

# Server Implementation

```
1 #include <arpa/inet.h>
2 #include <cstring>
3 #include <iostream>
4 #include <netinet/in.h>
5 #include <sys/socket.h>
6 #include <unistd.h>
7
8 using namespace std;
9
10 int main()
11 {
12     int serverSocket = socket(AF_INET, SOCK_STREAM, 0);
13
14     sockaddr_in serverAddress;
15     memset(&serverAddress, 0, sizeof(serverAddress));
16     serverAddress.sin_family = AF_INET;
17     serverAddress.sin_port = htons(8080);
18     serverAddress.sin_addr.s_addr = INADDR_ANY;
19
20     if (::bind(serverSocket,
21                (struct sockaddr*)&serverAddress,
22                sizeof(serverAddress)) < 0)
23     {
24         perror("Bind Failed");
25         return 1;
26     }
27 }
```

```
28     // Retrieve the bound IP/port
29     struct sockaddr_in bound_addr;
30     socklen_t len = sizeof(bound_addr);
31
32     if (getsockname(serverSocket, (struct sockaddr*)&bound_addr, &len) == 0)
33     {
34         char ip_str[INET_ADDRSTRLEN];
35         inet_ntop(AF_INET, &bound_addr.sin_addr, ip_str, sizeof(ip_str));
36
37         cout << "Server Socket bounded to IP " << ip_str
38             << " port " << ntohs(bound_addr.sin_port) << endl;
39     }
40     else
41     {
42         perror("getsockname failed");
43     }
44
45     listen(serverSocket, 5);
46     cout << "Server is now Listening" << endl;
47
48     int clientSocket = accept(serverSocket, nullptr, nullptr);
49
50     char buffer[1024] = {0};
51     recv(clientSocket, buffer, sizeof(buffer), 0);
52     cout << "Message from Client: " << buffer << endl;
53
54     close(serverSocket);
55     close(clientSocket);
56
57     return 0;
58 }
```

# Client Implementation

```
1 #include <iostream>
2 #include <cstring>
3 #include <sys/socket.h>
4 #include <arpa/inet.h>
5 #include <unistd.h>
6
7 int main(int argc, char *argv[]) {
8     if (argc < 2) {
9         std::cerr << "Usage: " << argv[0] << " <server_ip>\n";
10    return 1;
11 }
12
13 const char* server_ip = argv[1];
14
15 // Create socket
16 int sock = socket(AF_INET, SOCK_STREAM, 0);
17 if (sock < 0) {
18     std::cerr << "Error creating socket.\n";
19     return 1;
20 }
21
22 // Setup server address struct
23 struct sockaddr_in server_addr;
24 std::memset(&server_addr, 0, sizeof(server_addr));
25
```

```
26     server_addr.sin_family = AF_INET;
27     server_addr.sin_port = htons(8080);
28
29     // Convert IP string → binary form
30     if (inet_pton(AF_INET, server_ip, &(server_addr.sin_addr)) <= 0) {
31         std::cerr << "Invalid IP address format.\n";
32         close(sock);
33         return 1;
34     }
35
36     // Connect to server
37     if (connect(sock, (struct sockaddr*)&server_addr, sizeof(server_addr)) < 0) {
38         std::cerr << "Connect failed.\n";
39         close(sock);
40         return 1;
41     }
42
43     // Message to send
44     const char* msg = "Hello from client!";
45     send(sock, msg, strlen(msg), 0);
46
47     // Close socket
48     close(sock);
49     return 0;
50 }
```

# Build Process

## Function Calls:

**all** - compiles all the programs concurrently

**server** - builds the server executable

**client** - builds the client executable

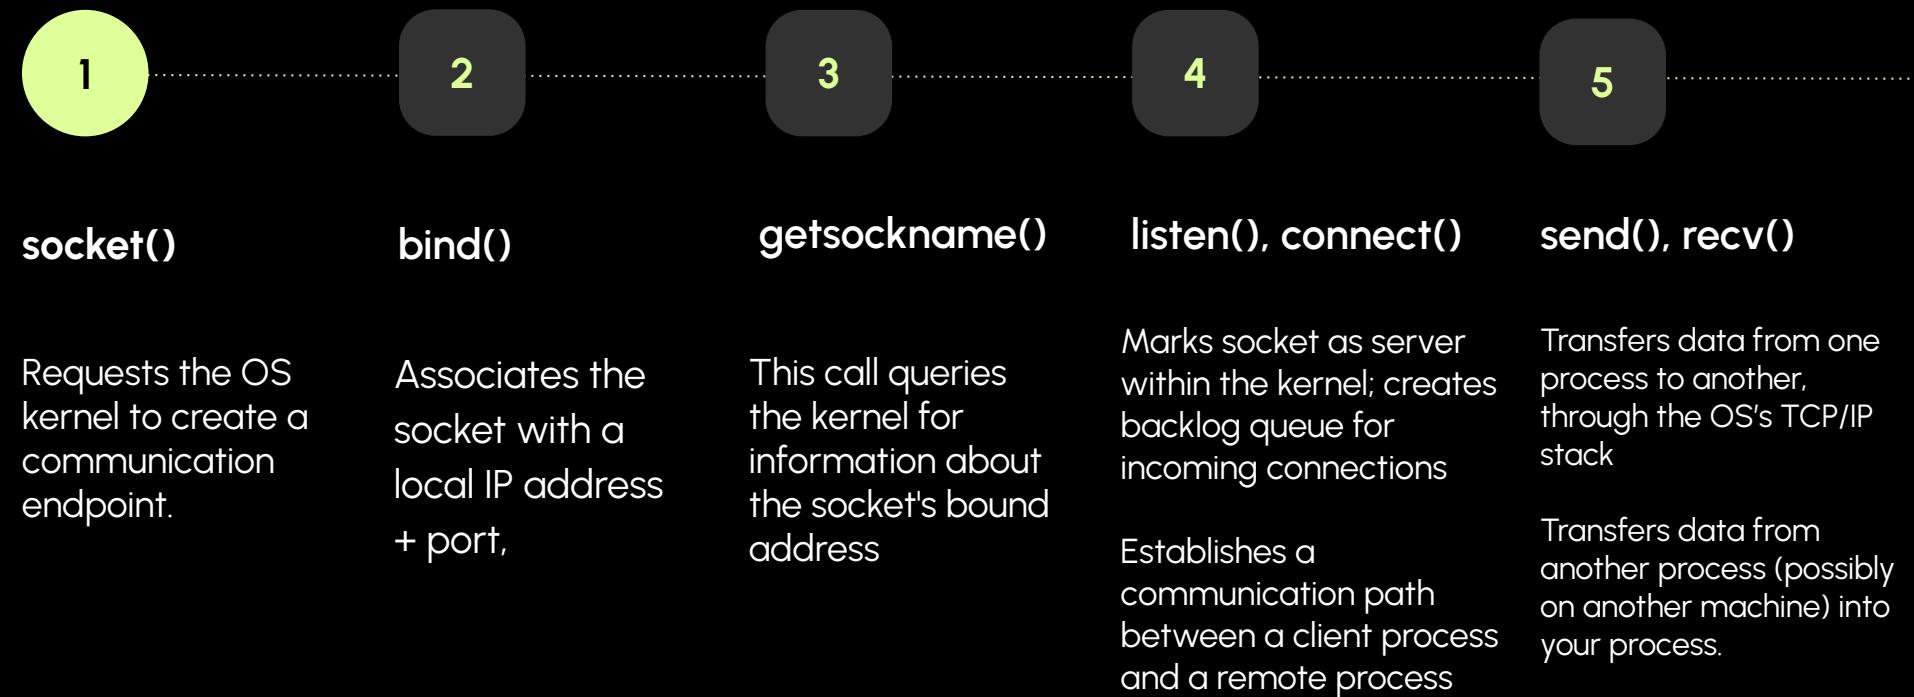
**clean** - removes compiled binaries

## Purpose:

- The Makefile automates compilation.
- Ensures consistent build commands.
- Streamlines process to run the program.

```
1 CXX = g++
2 CXXFLAGS = -Wall -Wextra
3
4 all: client server
5
6
7 client: client.cpp
8     $(CXX) $(CXXFLAGS) -o client client.cpp
9
10 server: server.cpp
11     $(CXX) $(CXXFLAGS) -o server server.cpp
12
13 run-client: client
14     ./client
15
16 run-server: server
17     ./server
18
19 clean:
20     rm -f client server
```

# OS IPC Breakdown



# Future Enhancements

We plan to implement the following improvements as we continue to build on this project:

## 1 Improved Interprocess Communication

As of now, we have implemented network-based message passing through TCP sockets, rather than using traditional pipes to convey message passing. Implementing pipe-based message passing, which would demonstrate communication across different address spaces

## 2 Add Multi-Client Support (Concurrency)

Right now, our server can only manage a single client connection at once. To improve scalability, we plan to implement multiprocessing so that each new client triggers a fork(). This will allow a separate child process to handle that client while the main server continues accepting others

## 3 Error Handling and Logging

Improved error handling and logging would make the system more robust and easier to maintain. Adding structured logs with timestamps and clear error messages would help diagnose issues more effectively. Implementing consistent error-recovery steps would bring the project closer to real-world network server reliability.

# Thank you

Questions ?